

Milestone 2

▼ Importing the libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report, confusion_matrix
import warnings
import pickle
from scipy import stats
warnings.filterwarnings('ignore')
plt.style.use('fivethirtyeight')
```

▼ Read the Dataset

```
data=pd.read_csv("/content/Copy of Data_Train.csv")
data.head()
```

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Dur
0	IndiGo	24/03/2019	Banglore	New Delhi	BLR ? DEL	22:20	01:10 22 Mar	2
1	Air India	1/05/2019	Kolkata	Banglore	CCU ? IXR ? BBI ? BLR	05:50	13:15	7

```
data.shape
```

```
(10683, 11)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10683 non-null  object
1   Date_of_Journey        10683 non-null  object
2   Source                 10683 non-null  object
3   Destination            10683 non-null  object
4   Route                  10682 non-null  object
5   Dep_Time               10683 non-null  object
6   Arrival_Time           10683 non-null  object
7   Duration               10683 non-null  object
8   Total_Stops            10682 non-null  object
9   Additional_Info        10683 non-null  object
10  Price                  10683 non-null  int64
dtypes: int64(1), object(10)
memory usage: 918.2+ KB
```

```
data.isnull().sum()
```

```
Airline                0
Date_of_Journey        0
Source                 0
Destination            0
Route                  1
Dep_Time               0
Arrival_Time           0
Duration               0
Total_Stops            1
Additional_Info        0
Price                  0
dtype: int64
```

▼ Data Preparation

```
category=['Airline','Source','Destination','Additional_Info']
category
```

```
['Airline', 'Source', 'Destination', 'Additional_Info']
```

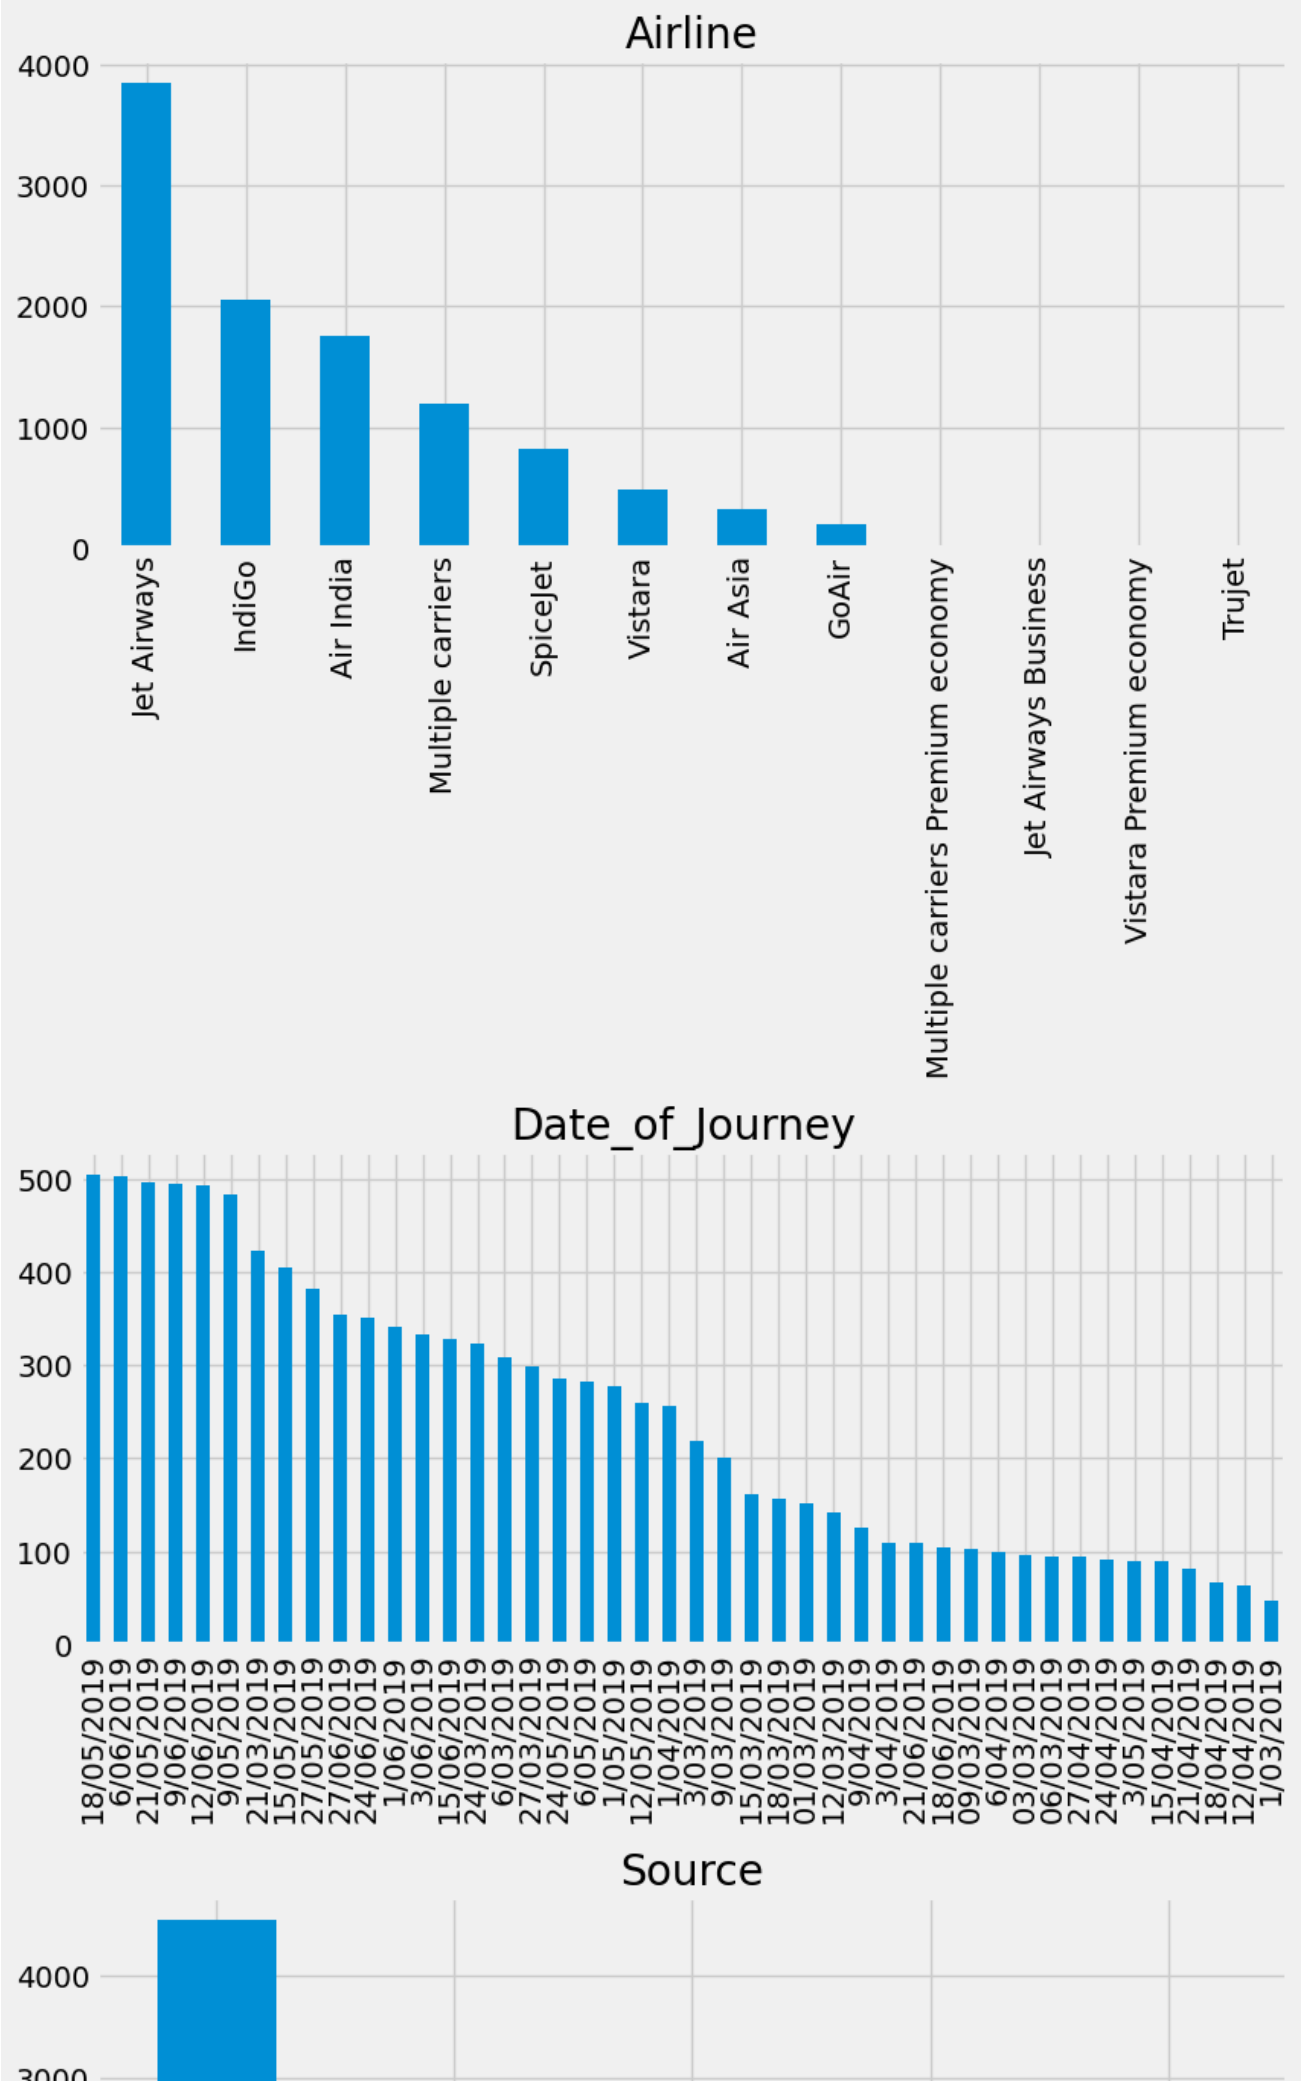
```
for i in category:
    print(i, data[i].unique())
```

```
Airline ['IndiGo' 'Air India' 'Jet Airways' 'SpiceJet' 'Multiple carriers' 'GoAir'
'Vistara' 'Air Asia' 'Vistara Premium economy' 'Jet Airways Business'
'Multiple carriers Premium economy' 'Trujet']
Source ['Bangalore' 'Kolkata' 'Delhi' 'Chennai' 'Mumbai']
Destination ['New Delhi' 'Bangalore' 'Cochin' 'Kolkata' 'Delhi' 'Hyderabad']
Additional_Info ['No info' 'In-flight meal not included' 'No check-in baggage include
'1 Short layover' 'No Info' '1 Long layover' 'Change airports'
'Business class' 'Red-eye flight' '2 Long layover']
```

```
category_cols=data.select_dtypes(include=['object']).columns
category_cols

Index(['Airline', 'Date_of_Journey', 'Source', 'Destination', 'Route',
      'Dep_Time', 'Arrival_Time', 'Duration', 'Total_Stops',
      'Additional_Info'],
      dtype='object')

#plotting a barchart for each of the categorical value
for column in category_cols:
    plt.figure(figsize=(20,4))
    plt.subplot(121)
    data[column].value_counts().plot(kind='bar')
    plt.title(column)
```



```
data.Route=data.Route.str.split('->')
```

```
data.Route
```

```
0          [BLR ? DEL]
1      [CCU ? IXR ? BBI ? BLR]
2      [DEL ? LKO ? BOM ? COK]
3          [CCU ? NAG ? BLR]
4          [BLR ? NAG ? DEL]
...
10678      [CCU ? BLR]
10679      [CCU ? BLR]
10680      [BLR ? DEL]
10681      [BLR ? DEL]
10682      [DEL ? GOI ? BOM ? COK]
Name: Route, Length: 10683, dtype: object
```

```
data['City1']=data.Route.str[0]
data['City2']=data.Route.str[1]
data['City3']=data.Route.str[2]
data['City4']=data.Route.str[3]
data['City5']=data.Route.str[4]
data['City6']=data.Route.str[5]
```

```
data.Date_of_Journey=data.Date_of_Journey.str.split('/')
```

```
data.Date_of_Journey
```

```
0          [24, 03, 2019]
1          [1, 05, 2019]
2          [9, 06, 2019]
3          [12, 05, 2019]
4          [01, 03, 2019]
...
10678      [9, 04, 2019]
10679      [27, 04, 2019]
10680      [27, 04, 2019]
10681      [01, 03, 2019]
10682      [9, 05, 2019]
Name: Date_of_Journey, Length: 10683, dtype: object
```

```
data['Date']=data.Date_of_Journey.str[0]
data['Month']=data.Date_of_Journey.str[1]
data['Year']=data.Date_of_Journey.str[2]
```

```
data.Dep_Time=data.Dep_Time.str.split(':')
```

```
data['Dep_Time_Hour']=data.Dep_Time.str[0]
data['Dep_Time_Mins']=data.Dep_Time.str[1]
```

```
data.Arrival_Time=data.Arrival_Time.str.split(' ')
```

```
data['Arrival_date']=data.Arrival_Time.str[1]
data['Time_of_Arrival']=data.Arrival_Time.str[0]
```

```
data['Time_of_Arrival']=data.Time_of_Arrival.str.split(':')
```

```
data['Arrival_Time_Hour']=data.Time_of_Arrival.str[0]
data['Arrival_Time_Mins']=data.Time_of_Arrival.str[1]
```

```
data.Duration=data.Duration.str.split(' ')
```

```
data['Travel_Hours']=data.Duration.str[0]
data['Travel_Hours']=data['Travel_Hours'].str.split('h')
data['Travel_Hours']=data['Travel_Hours'].str[0]
data.Travel_Hours=data.Travel_Hours
data['Travel_Mins']=data.Duration.str[1]
```

```
data.Travel_Mins=data.Travel_Mins.str.split('m')
data.Travel_Mins=data.Travel_Mins.str[0]
```

```
data.Total_Stops.replace('non_stop',0,inplace=True)
data.Total_Stops=data.Total_Stops.str.split(' ')
data.Total_Stops=data.Total_Stops.str[0]
```

```
data.Additional_Info.unique()
```

```
array(['No info', 'In-flight meal not included',
       'No check-in baggage included', '1 Short layover', 'No Info',
       '1 Long layover', 'Change airports', 'Business class',
       'Red-eye flight', '2 Long layover'], dtype=object)
```

```
data.Additional_Info.replace('No Info','No info',inplace=True)
```

```
data.isnull().sum()
```

```
Airline          0
Date_of_Journey  0
Source           0
Destination      0
Route           1
Dep_Time         0
Arrival_Time     0
Duration         0
Total_Stops      1
Additional_Info   0
Price            0
City1            1
City2          10683
City3          10683
```

```

City4          10683
City5          10683
City6          10683
Date           0
Month          0
Year           0
Dep_Time_Hour  0
Dep_Time_Mins  0
Arrival_date   6348
Time_of_Arrival 0
Arrival_Time_Hour 0
Arrival_Time_Mins 0
Travel_Hours   0
Travel_Mins    1032
dtype: int64

```

```
data.drop(['City4', 'City5', 'City6'], axis=1, inplace=True)
```

```
data.drop(['Date_of_Journey', 'Route', 'Dep_Time', 'Arrival_Time', 'Duration'], axis=1, inplace=True)
data.drop(['Time_of_Arrival'], axis=1, inplace=True)
```

▼ Replacing Missing Values

```
data.isnull().sum()
```

```

Airline          0
Source           0
Destination       0
Total_Stops      1
Additional_Info   0
Price            0
City1            1
City2           10683
City3           10683
Date             0
Month            0
Year             0
Dep_Time_Hour    0
Dep_Time_Mins    0
Arrival_date     6348
Arrival_Time_Hour 0
Arrival_Time_Mins 0
Travel_Hours     0
Travel_Mins      1032
dtype: int64

```

```
data['City3'].fillna('None', inplace=True)
```

```
data['Arrival_date'].fillna(data['Date'], inplace=True)
```

```
data['Travel_Mins'].fillna(0, inplace=True)
```

```
#data.Total_Stops=data.Total_Stops.astype('int64')
data.Date=data.Date.astype('int64')
data.Month=data.Month.astype('int64')
data.Year=data.Year.astype('int64')
data.Dep_Time_Hour=data.Dep_Time_Hour.astype('int64')
data.Dep_Time_Hour=data.Dep_Time_Hour.astype('int64')
data.Dep_Time_Mins=data.Dep_Time_Mins.astype('int64')
data.Arrival_date=data.Arrival_date.astype('int64')
data.Arrival_Time_Hour=data.Arrival_Time_Hour.astype('int64')
data.Arrival_Time_Mins=data.Arrival_Time_Mins.astype('int64')
#data.Travel_Hours=data.Travel_Hours.astype('int64')
data.Travel_Mins=data.Travel_Mins.astype('int64')
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10683 non-null  object
1   Source                 10683 non-null  object
2   Destination            10683 non-null  object
3   Total_Stops            10682 non-null  object
4   Additional_Info        10683 non-null  object
5   Price                  10683 non-null  int64
6   City1                  10682 non-null  object
7   City2                  0 non-null      float64
8   City3                  10683 non-null  object
9   Date                   10683 non-null  int64
10  Month                  10683 non-null  int64
11  Year                   10683 non-null  int64
12  Dep_Time_Hour          10683 non-null  int64
13  Dep_Time_Mins          10683 non-null  int64
14  Arrival_date           10683 non-null  int64
15  Arrival_Time_Hour      10683 non-null  int64
16  Arrival_Time_Mins      10683 non-null  int64
17  Travel_Hours           10683 non-null  object
18  Travel_Mins            10683 non-null  int64
dtypes: float64(1), int64(10), object(8)
memory usage: 1.5+ MB
```

```
data[data['Travel_Hours']=='5m']
```


Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2
---------	--------	-------------	-------------	-----------------	-------	-------	-------

```
data.drop(index=6474,inplace=True,axis=0)
```

6474	Air India	Mumbai	Hyderabad	2	No info	17227	HYD
------	-----------	--------	-----------	---	---------	-------	-----

```
data.Travel_Hours=data.Travel_Hours.astype('int64')
```

HYD

```
categorical=['Airline','Source','Destination','Additional_Info','City1']
```

```
numerical=['Total_Stops','Date','Month','Year','Dep_Time_Hour','Dep_Time_Mins','Arrival_da
```

▼ visual Analysis

▼ sample visualization

```
import numpy as np
from matplotlib import pyplot as plt
```

```
ys=200+np.random.randn(100)
x=[x for x in range(len(ys))]
```

```
plt.plot(x,ys,'-')
plt.fill_between(x,ys,195,where=(ys>195),facecolor='g',alpha=0.6)
```

```
plt.title("sample visualization")
plt.show()
```

203

sample visualization

```
#plotting countplots for categorical data
```

```
import seaborn as sns
```

```
c=1
```

```
plt.figure(figsize=(20,45))
```

```
for i in categorical:
```

```
    plt.subplot(6,3,c)
```


```
    sns.countplot(data['Price'])
```

```
    plt.xticks(rotation=90)
```


```
    plt.tight_layout(pad=3.0)
```

```
    c=c+1
```

```
plt.show()
```




▼ Distribution of price column



```
plt.figure(figsize=(15,8))  
sns.distplot(data.Price)
```

<Axes: xlabel='Price', ylabel='Density'>



data.columns

```
Index(['Airline', 'Source', 'Destination', 'Total_Stops', 'Additional_Info',
      'Price', 'City1', 'City2', 'City3', 'Date', 'Month', 'Year',
      'Dep_Time_Hour', 'Dep_Time_Mins', 'Arrival_date', 'Arrival_Time_Hour',
      'Arrival_Time_Mins', 'Travel_Hours', 'Travel_Mins'],
      dtype='object')
```

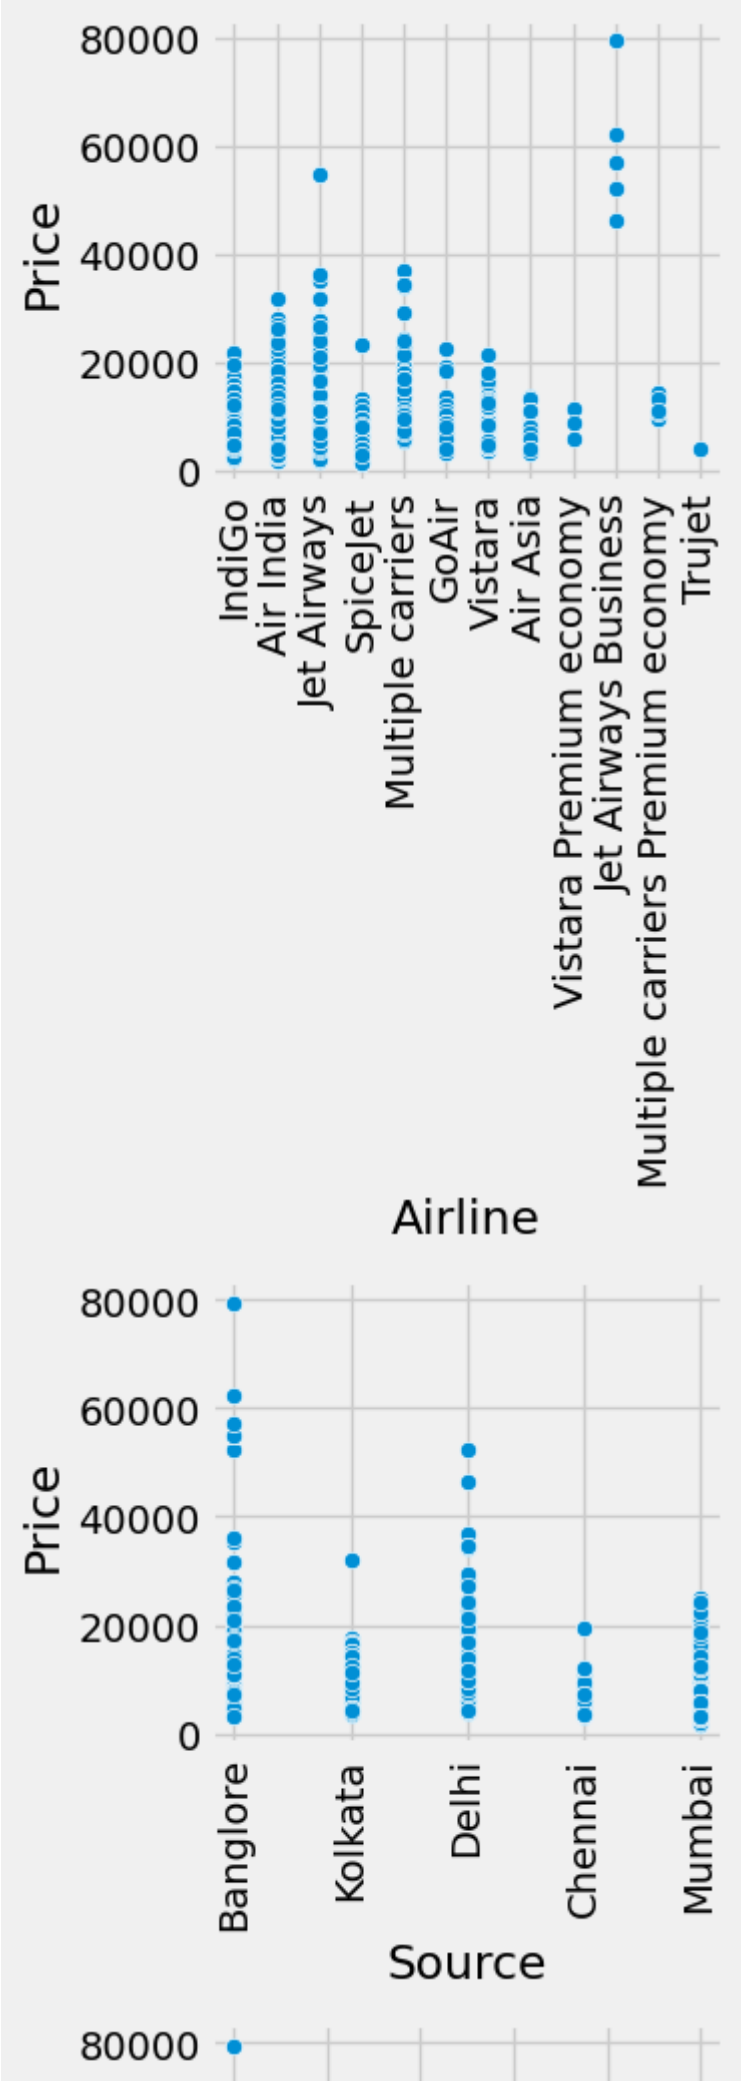
Checking the relation of price with categorical data

```
import seaborn as sns
c=1
```

```
for i in categorical:
    plt.figure(figsize=(10,20))

    plt.subplot(6,3,c)

    sns.scatterplot(x=data[i],y=data.Price)
    plt.xticks(rotation=90)
    #plt.tight_layout(pad=3.0)
    c=c+1
    plt.show()
```



```
# Checking flight with high prices
data[data.Price>50000]
data.head()
pd.set_option('display.max_columns',25)
data.head()
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2
0	IndiGo	Banglore	New Delhi	non-stop	No info	3897	BLR ? DEL	NaN
1	Air India	Kolkata	Banglore	2	No info	7662	CCU ? IXR ? BBI ? BLR	NaN
2	Jet Airways	Delhi	Cochin	2	No info	13882	DEL ? LKO ? BOM ? COK	NaN
3	IndiGo	Kolkata	Banglore	1	No info	6218	CCU ? NAG ? BLR	NaN
4	IndiGo	Banglore	New Delhi	1	No info	13302	BLR ? NAG ? DEL	NaN



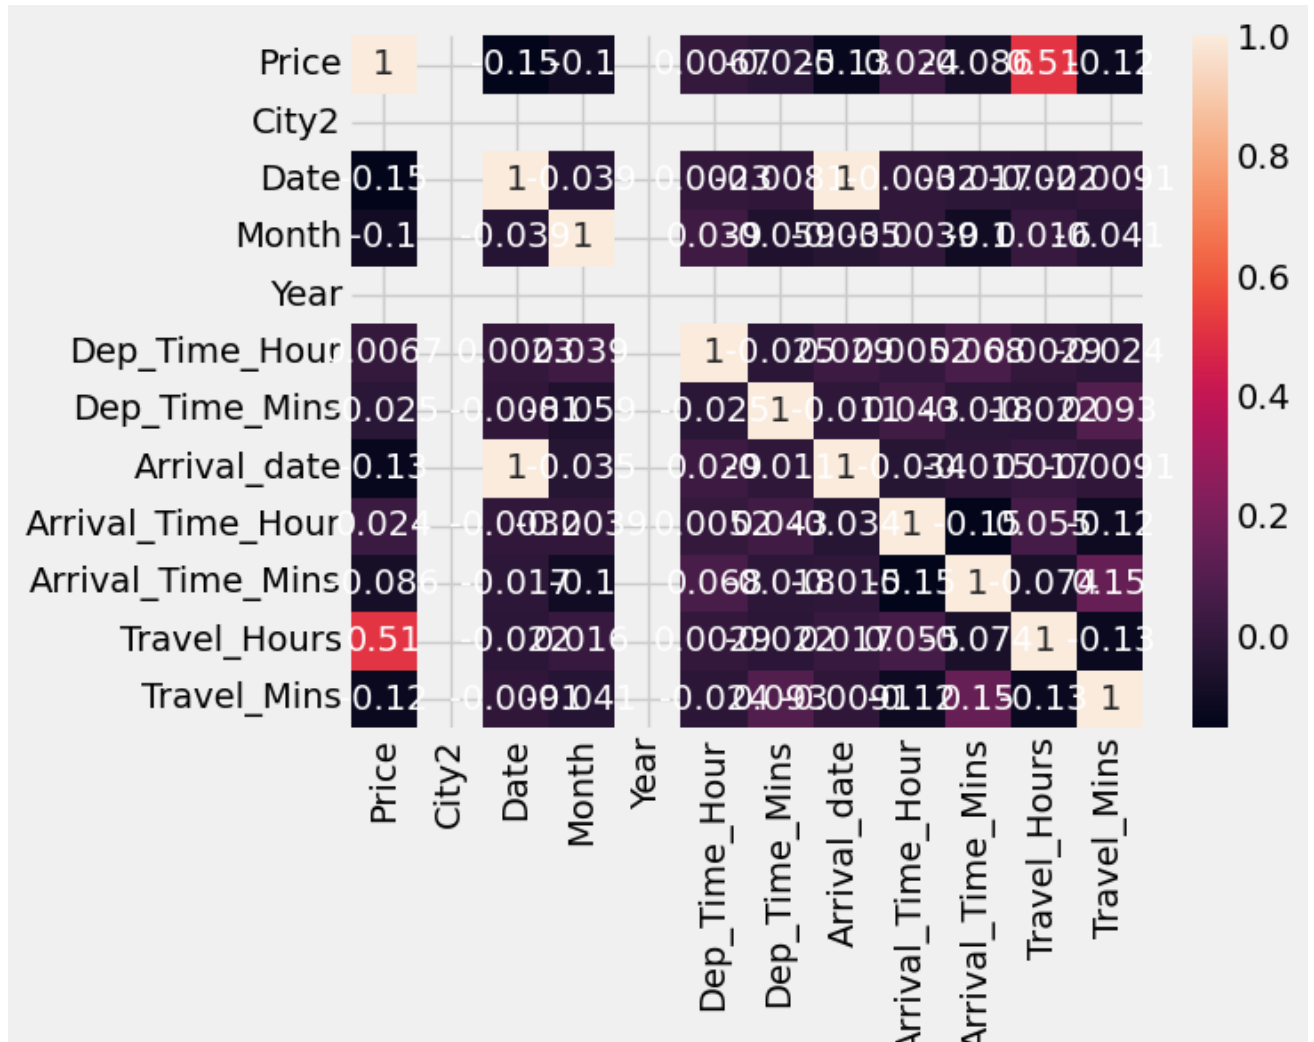
```
data['Year'].max()
```

2019

▼ Checking the correlation using HeatMap

```
sns.heatmap(data.corr(),annot=True)
```

<Axes: >



data.info()


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10682 entries, 0 to 10682
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Airline                10682 non-null  object
1   Source                 10682 non-null  object
2   Destination            10682 non-null  object
3   Total_Stops            10681 non-null  object
4   Additional_Info        10682 non-null  object
5   Price                  10682 non-null  int64
6   City1                  10681 non-null  object
7   City2                  0 non-null      float64
8   City3                  10682 non-null  object
9   Date                   10682 non-null  int64
10  Month                  10682 non-null  int64
11  Year                   10682 non-null  int64
12  Dep_Time_Hour          10682 non-null  int64
13  Dep_Time_Mins          10682 non-null  int64
14  Arrival_date           10682 non-null  int64
15  Arrival_Time_Hour      10682 non-null  int64
16  Arrival_Time_Mins      10682 non-null  int64
17  Travel_Hours            10682 non-null  int64
18  Travel_Mins            10682 non-null  int64
dtypes: float64(1), int64(11), object(7)
memory usage: 1.9+ MB
```

data

Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2
---------	--------	-------------	-------------	-----------------	-------	-------	-------

```
# Checking relation price with numerical values  
c=1
```


```
for i in numerical:  
    plt.figure(figsize=(10,20))  
    plt.subplot(6,3,c)  
    sns.scatterplot(x=data[i],y=data.Price)  
    plt.xticks(rotation=90)  
    #plt.tight_layout(pad=3.0)  
    c=c+1  
    plt.show()
```



▼ Outlier detection for Price column

```
# Detecting the outliers
import seaborn as sns
sns.boxplot(data['Price'])
```

<Axes: >



▼ Lable Encodig

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```

```
data.Airline=le.fit_transform(data.Airline)
data.Source=le.fit_transform(data.Source)
data.Destination=le.fit_transform(data.Destination)
data.Total_Stops=le.fit_transform(data.Total_Stops)
data.City1=le.fit_transform(data.City1)
data.City2=le.fit_transform(data.City2)
data.City3=le.fit_transform(data.City3)
data.Additional_Info=le.fit_transform(data.Additional_Info)
data.head()
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2	C
0	3	0	5	4	7	3897	18	0	
1	1	3	0	1	7	7662	84	0	
2	4	2	1	1	7	13882	118	0	
3	3	3	0	0	7	6218	91	0	
4	3	0	5	0	7	13302	29	0	



▼ Output Columns




```
data.head()
```

	Airline	Source	Destination	Total_Stops	Additional_Info	Price	City1	City2	C
0	3	0	5	4	7	3897	18	0	
1	1	3	0	1	7	7662	84	0	
2	4	2	1	1	7	13882	118	0	
3	3	3	0	0	7	6218	91	0	
4	3	0	5	0	7	13302	29	0	



```
data=data[['Airline','Source','Destination','Date','Month','Year','Dep_Time_Hour','Dep_Tim
```



```
data.head()
```

	Airline	Source	Destination	Date	Month	Year	Dep_Time_Hour	Dep_Time_Mins	Arr
0	3	0	5	24	3	2019	22	20	
1	1	3	0	1	5	2019	5	50	
2	4	2	1	9	6	2019	9	25	
3	3	3	0	12	5	2019	18	5	
4	3	0	5	1	3	2019	16	50	



Milestone 3

Travel Hours

▼ Exploratory Data Analysis

▼ Descriptive Statistical

```
data.describe()
```

	Airline	Source	Destination	Date	Month	Year	D
count	10682.000000	10682.000000	10682.000000	10682.000000	10682.000000	10682.0	
mean	3.966205	1.952069	1.435967	13.509081	4.708762	2019.0	
std	2.352090	1.177110	1.474773	8.479363	1.164294	0.0	
min	0.000000	0.000000	0.000000	1.000000	3.000000	2019.0	
25%	3.000000	2.000000	0.000000	6.000000	3.000000	2019.0	
50%	4.000000	2.000000	1.000000	12.000000	5.000000	2019.0	
75%	4.000000	3.000000	2.000000	21.000000	6.000000	2019.0	
max	11.000000	4.000000	5.000000	27.000000	6.000000	2019.0	

▼ Scaling the Data

```
from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
```

```
data1=ss.fit_transform(data)
```

```
data1=pd.DataFrame(data1,columns=data.columns)
data1.head()
```

	Airline	Source	Destination	Date	Month	Year	Dep_Time_Hour	Dep_Tim
0	-0.410805	-1.658435	2.416778	1.237288	-1.467707	0.0	1.654268	-0.0
1	-1.261152	0.890299	-0.973732	-1.475307	0.250153	0.0	-1.303000	1.0
2	0.014369	0.040721	-0.295630	-0.531796	1.109082	0.0	-0.607172	0.0
3	-0.410805	0.890299	-0.973732	-0.177979	0.250153	0.0	0.958440	-1.0
4	-0.410805	-1.658435	2.416778	-1.475307	-1.467707	0.0	0.610527	1.0

```
y=data1['Price']
x=data1.drop(columns=['Price'],axis=1)
```

▼ Splitting data into train and test

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
x_train.head()
```

	Airline	Source	Destination	Date	Month	Year	Dep_Time_Hour	Dep_
10004	0.864716	0.040721	-0.29563	1.591104	0.250153	0.0	-0.781129	
3684	0.014369	0.040721	-0.29563	-0.531796	0.250153	0.0	-0.259258	
1034	1.715063	0.040721	-0.29563	1.237288	-0.608777	0.0	0.436570	
3909	0.864716	0.040721	-0.29563	0.883471	-1.467707	0.0	-0.085301	
3088	-1.261152	0.040721	-0.29563	1.237288	1.109082	0.0	0.784483	

```
x_train.shape
```

```
(8545, 9)
```

```
y_train.shape
```

```
(8545,)
```

Milestone 4

▼ Model building

▼ Using Ensemble Techniques

```
from sklearn.ensemble import RandomForestRegressor,GradientBoostingRegressor,AdaBoostRegressor
rfr=RandomForestRegressor()
gb=GradientBoostingRegressor()
ad=AdaBoostRegressor()
```

```
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```

for i in [rfr,gb,ad]:
    i.fit(x_train,y_train)
    y_pred=i.predict(x_test)
    test_score=r2_score(y_test,y_pred)
    train_score=r2_score(y_train,i.predict(x_train))
    if abs(train_score-test_score)<=0.2:
        print(i)

    print("R2 score is",r2_score(y_test,y_pred))
    print("R2 for train data",r2_score(y_train,i.predict(x_train)))
    print("Mean Absolute Error is",mean_absolute_error(y_pred,y_test))
    print("Mean Squared Error is",mean_squared_error(y_pred,y_test))
    print("Root Mean Squared Error is",(mean_squared_error(y_pred,y_test,squared=False

```

```

RandomForestRegressor()
R2 score is 0.833020601170679
R2 for train data 0.9105955758611345
Mean Absolute Error is 0.27361502073943944
Mean Squared Error is 0.16604267786110014
Root Mean Squared Error is 0.4074833467285505
GradientBoostingRegressor()
R2 score is 0.7586150253290672
R2 for train data 0.7199250828810017
Mean Absolute Error is 0.3700514763508659
Mean Squared Error is 0.24003085333157612
Root Mean Squared Error is 0.4899294370943392
AdaBoostRegressor()
R2 score is 0.3275715393439753
R2 for train data 0.348524998044629
Mean Absolute Error is 0.6374032115068605
Mean Squared Error is 0.6686562717324748
Root Mean Squared Error is 0.817714052546778

```

▼ Regression Model

```

from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor

```

```

from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error

```

```

knn=KNeighborsRegressor()
svr=SVR()
dt=DecisionTreeRegressor()

for i in [knn,svr,dt]:
    i.fit(x_train,y_train)
    y_pred=i.predict(x_test)
    test_score=r2_score(y_test,y_pred)
    train_score=r2_score(y_train,i.predict(x_train))
    if abs(train_score-test_score)<=0.1:
        print(i)
        print('R2 Score is',r2_score(y_test,y_pred))
        print('R2 Score for train data',r2_score(y_train,i.predict(x_train)))
        print('Mean Absolute Error is',mean_absolute_error(y_test,y_pred))
        print('Mean Squared Error is',mean_squared_error(y_test,y_pred))
        print('Root Mean Squared Error is',(mean_squared_error(y_test,y_pred,squared=False)))

KNeighborsRegressor()
R2 Score is 0.7067031916509059
R2 Score for train data 0.7708199280026578
Mean Absolute Error is 0.36928172590251457
Mean Squared Error is 0.2916514720248588
Root Mean Squared Error is 0.5400476571793075
SVR()
R2 Score is 0.5890645626885227
R2 Score for train data 0.5570926824460594
Mean Absolute Error is 0.44723971869314205
Mean Squared Error is 0.40863017185110667
Root Mean Squared Error is 0.6392418727298038

```

▼ Checking cross validation for RandomForestRegressor

```

from sklearn.model_selection import cross_val_score
for i in range(2,5):
    cv=cross_val_score(rfr,x,y,cv=i)
    print(rfr,cv.mean())

RandomForestRegressor() 0.7687854467304049
RandomForestRegressor() 0.7696443473741302
RandomForestRegressor() 0.7817983669249348

```

rfr

▼ **RandomForestRegressor**


RandomForestRegressor(max_features='sqrt', n_estimators=10)

▼ Hypertuning the model

```
from sklearn.model_selection import RandomizedSearchCV
```


```
param_grid={'n_estimators':[10,30,50,70,100],'max_depth':[None,1,2,3],'max_features':['auto', 'sqrt', 'log2']}
rfr=RandomForestRegressor()
rf_res=RandomizedSearchCV(estimator=rfr,param_distributions=param_grid,cv=3,verbose=2,n_jobs=-1)
rf_res.fit(x_train,y_train)
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits



```
gb=GradientBoostingRegressor()
gb_res=RandomizedSearchCV(estimator=gb,param_distributions=param_grid,cv=3,verbose=2,n_jobs=-1)
gb_res.fit(x_train,y_train)
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits



▼ Accuracy

```
rfr=RandomForestRegressor(n_estimators=10,max_features='sqrt',max_depth=None)
rfr.fit(x_train,y_train)
y_train_pred=rfr.predict(x_train)
y_test_pred=rfr.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.8830371391616789
test accuracy 0.7445652199866283
```

▼ Checking train and test accuracy by RandomSearchCV using KNN model2

```
knn=KNeighborsRegressor(n_neighbors=2,algorithm='auto',metric_params=None,n_jobs=-1)
knn.fit(x_train,y_train)
```

```
y_train_pred=knn.predict(x_train)
y_test_pred=knn.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.8252359370660914
test accuracy 0.6531487599455481
```

▼ Checking cross validation for RandomForestRegressor

```
from sklearn.model_selection import cross_val_score
for i in range(2,5):
    cv=cross_val_score(rfr,x,y,cv=i)
    print(rfr,cv.mean())
```

```
RandomForestRegressor() 0.7655573723652831
RandomForestRegressor() 0.7724983811643485
RandomForestRegressor() 0.7836926794833963
```

```
rfr
```

```
▼ RandomForestRegressor
RandomForestRegressor(max_features='sqrt', n_estimators=10)
```

▼ Evaluating performance of the model and saving the model

```
rfr=RandomForestRegressor(n_estimators=10,max_features='sqrt',max_depth=None)
rfr.fit(x_train,y_train)
y_train_pred=rfr.predict(x_train)
y_test_pred=rfr.predict(x_test)
print("train accuracy",r2_score(y_train_pred,y_train))
print("test accuracy",r2_score(y_test_pred,y_test))
```

```
train accuracy 0.8850928588419827
test accuracy 0.7702120755570521
```

```
Predicted_values=pd.DataFrame({'Actual':y_test,'Predicted':y_pred})
```

```
Predicted_values
```

	Actual	Predicted
6075	1.641563	1.681688
3544	-0.895161	-0.895161
9290	0.021842	-0.110966
5032	-1.133955	-1.190563
2483	0.826714	1.171675
...
9796	-0.364002	0.824871
9870	-0.968253	-0.614942

```
Prices=rfr.predict(x_test)

price_list=pd.DataFrame({'Price':Prices})
```

price_list

	Price
0	0.902105
1	-0.764443
2	-0.032446
3	-1.168961
4	0.900212
...	...
2132	0.710256
2133	-0.938430
2134	-0.368481
2135	0.234306
2136	0.670456

2137 rows × 1 columns

▼ Milestone 6

```
import pickle
pickle.dump(rfr,open('model1.pk1','wb'))
```

✓ 0s completed at 11:50 AM

● ✕